

IMPORTING REQUIRED LIBRARIES

```
In [1]: import pandas as ps
import numpy as ns
import seaborn as sn
import matplotlib.pyplot as pl
import warnings
warnings.filterwarnings('ignore')
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

In [2]: df=ps.read_csv(r"C:\Users\Samdure\OneDrive\Desktop\American_Housing_Data.csv")
```

DATA CLEANING

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 39981 entries, 0 to 39980
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Zip Code                              39981 non-null  int64
1   Price                                39981 non-null  float64
2   Beds                                 39981 non-null  int64
3   Baths                               39981 non-null  int64
4   Living Space                         39981 non-null  int64
5   Address                             39981 non-null  object
6   City                                39981 non-null  object
7   State                               39981 non-null  object
8   Zip Code Population                  39981 non-null  int64
9   Zip Code Density                    39981 non-null  float64
10  County                              39981 non-null  object
11  Median Household Income              39979 non-null  float64
12  Latitude                             39981 non-null  float64
13  Longitude                            39981 non-null  float64
dtypes: float64(5), int64(5), object(4)
memory usage: 4.3+ MB

In [4]: df.head(10)
```

Out[4]:

| | Zip Code | Price | Beds | Baths | Living Space | Address | City | State | Zip Code Population | Zip Code Density | County | Median Household Income | Latitude | Longitude |
|---|----------|-----------|------|-------|--------------|------------------------|----------|----------|---------------------|------------------|----------|-------------------------|----------|-----------|
| 0 | 10013 | 3999000.0 | 2 | 3 | 1967 | 74 GRAND ST APT 3 | New York | New York | 29563 | 20967.9 | New York | 370046.0 | 40.72001 | -74.00472 |
| 1 | 10013 | 3999000.0 | 2 | 3 | 1967 | 74 GRAND ST APT 3 | New York | New York | 29563 | 20967.9 | New York | 370046.0 | 40.72001 | -74.00472 |
| 2 | 10014 | 1650000.0 | 1 | 1 | 718 | 140 CHARLES ST APT 4D | New York | New York | 29815 | 23740.9 | New York | 249880.0 | 40.73407 | -74.00601 |
| 3 | 10014 | 760000.0 | 3 | 2 | 1538 | 38 JONES ST | New York | New York | 29815 | 23740.9 | New York | 249880.0 | 40.73407 | -74.00601 |
| 4 | 10014 | 1100000.0 | 1 | 1 | 600 | 81 BEDFORD ST APT 3F | New York | New York | 29815 | 23740.9 | New York | 249880.0 | 40.73407 | -74.00601 |
| 5 | 10017 | 764900.0 | 1 | 1 | 643 | 145 E 48TH ST APT 11E | New York | New York | 15514 | 20107.7 | New York | 188289.0 | 40.75235 | -73.97260 |
| 6 | 10021 | 2499000.0 | 2 | 2 | 1471 | 234 E 70TH ST APT 4 | New York | New York | 42484 | 46004.0 | New York | 261254.0 | 40.76963 | -73.95899 |
| 7 | 10022 | 4580000.0 | 2 | 3 | 1800 | 641 5TH AVE # 29D | New York | New York | 33303 | 28998.9 | New York | 281977.0 | 40.75856 | -73.96787 |
| 8 | 10026 | 540000.0 | 2 | 1 | 750 | 45 CENTRAL PARK N # 4D | New York | New York | 39401 | 39689.7 | New York | 117438.0 | 40.80302 | -73.95348 |
| 9 | 10026 | 570000.0 | 1 | 1 | 589 | 300 W 110TH ST APT 19H | New York | New York | 39401 | 39689.7 | New York | 117438.0 | 40.80302 | -73.95348 |

```
In [5]: df.tail(5)
```

Out[5]:

| | Zip Code | Price | Beds | Baths | Living Space | Address | City | State | Zip Code Population | Zip Code Density | County | Median Household Income | Latitude | Longitude |
|-------|----------|-----------|------|-------|--------------|------------------------|---------|------------|---------------------|------------------|--------|-------------------------|----------|------------|
| 39976 | 98199 | 2495000.0 | 4 | 4 | 3380 | 2626 27TH AVE W | Seattle | Washington | 22890 | 2086.8 | King | 205611.0 | 47.65139 | -122.40223 |
| 39977 | 98199 | 2295000.0 | 4 | 4 | 2878 | 3215 32ND AVE W | Seattle | Washington | 22890 | 2086.8 | King | 205611.0 | 47.65139 | -122.40223 |
| 39978 | 98199 | 950000.0 | 3 | 2 | 1380 | 3257 22ND AVE W | Seattle | Washington | 22890 | 2086.8 | King | 205611.0 | 47.65139 | -122.40223 |
| 39979 | 98199 | 425000.0 | 2 | 1 | 856 | 3711 26TH PL W APT 102 | Seattle | Washington | 22890 | 2086.8 | King | 205611.0 | 47.65139 | -122.40223 |
| 39980 | 98199 | 1150000.0 | 3 | 3 | 2840 | 2911 25TH AVE W | Seattle | Washington | 22890 | 2086.8 | King | 205611.0 | 47.65139 | -122.40223 |

```
In [6]: df.isnull()
```

Out[6]:

| | Zip Code | Price | Beds | Baths | Living Space | Address | City | State | Zip Code Population | Zip Code Density | County | Median Household Income | Latitude | Longitude |
|-------|----------|-------|-------|-------|--------------|---------|-------|-------|---------------------|------------------|--------|-------------------------|----------|-----------|
| 0 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 1 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 2 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 3 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 4 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 39976 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 39977 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 39978 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 39979 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |
| 39980 | False | False | False | False | False | False | False | False | False | False | False | False | False | False |

39981 rows × 14 columns

```
In [7]: df.describe()
```

Out[7]:

| | Zip Code | Price | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|--------------|--------------|--------------|--------------|--------------|---------------------|------------------|-------------------------|--------------|--------------|
| count | 39981.000000 | 3.998100e+04 | 39981.000000 | 39981.000000 | 39981.000000 | 39981.000000 | 39981.000000 | 39979.000000 | 39981.000000 | 39981.000000 |
| mean | 64833.391336 | 6.227771e+05 | 3.171682 | 2.466572 | 1901.522723 | 37726.201996 | 2379.412483 | 110837.259861 | 36.435668 | -98.080576 |
| std | 25614.601116 | 9.469793e+05 | 1.308796 | 1.323042 | 1211.307257 | 18672.647445 | 2946.574792 | 47309.055715 | 4.446862 | 15.061145 |
| min | 10013.000000 | 1.800000e+03 | 1.000000 | 1.000000 | 2.000000 | 0.000000 | 0.000000 | 27475.000000 | 25.729830 | -122.826870 |
| 25% | 40215.000000 | 2.650000e+05 | 3.000000 | 2.000000 | 1200.000000 | 24465.000000 | 902.400000 | 76640.000000 | 33.239850 | -111.636310 |
| 50% | 74136.000000 | 3.999000e+05 | 3.000000 | 2.000000 | 1639.000000 | 35049.000000 | 1588.700000 | 100405.000000 | 36.166620 | -96.839680 |
| 75% | 85730.000000 | 6.749900e+05 | 4.000000 | 3.000000 | 2265.000000 | 46816.000000 | 2736.800000 | 135075.000000 | 39.283090 | -85.656980 |
| max | 98199.000000 | 3.800000e+07 | 54.000000 | 66.000000 | 74340.000000 | 116469.000000 | 58289.600000 | 900203.000000 | 47.742370 | -73.704510 |

```
In [8]: df.shape
```

Out[8]: (39981, 14)

```
In [9]: df.isnull().sum()
```

Out[9]: Zip Code 0
Price 0
Beds 0
Baths 0
Living Space 0
Address 0
City 0
State 0
Zip Code Population 0
Zip Code Density 0
County 0
Median Household Income 2
Latitude 0
Longitude 0
dtype: int64

In [10]:

df.columns

Out[10]:

Index(['Zip Code', 'Price', 'Beds', 'Baths', 'Living Space', 'Address', 'City',
'State', 'Zip Code Population', 'Zip Code Density', 'County',
'Median Household Income', 'Latitude', 'Longitude'],
dtype='object')

In [11]:

df.sample(4)

Out[11]:

| | Zip Code | Price | Beds | Baths | Living Space | Address | City | State | Zip Code Population | Zip Code Density | County | Median Household Income | Latitude | Longitude | |
|--|----------|-------|----------|-------|--------------|---------|--------------------------------|---------------|---------------------|------------------|--------|-------------------------|----------|-----------|------------|
| | 87 | 10306 | 798000.0 | 3 | 2 | 1750 | 305 HUSSON ST | Staten Island | New York | 55805 | 2980.6 | Richmond | 118424.0 | 40.57149 | -74.12430 |
| | 10927 | 43207 | 164900.0 | 3 | 1 | 924 | 2943 PARSONS AVE | Columbus | Ohio | 46038 | 768.8 | Franklin | 67173.0 | 39.89599 | -82.96392 |
| | 33837 | 92107 | 595000.0 | 2 | 2 | 986 | 4444 W POINT LOMA BLVD UNIT 37 | San Diego | California | 29753 | 3798.5 | San Diego | 133969.0 | 32.74020 | -117.24357 |
| | 807 | 19125 | 699995.0 | 3 | 3 | 2406 | 3000 CITYVIEW WALK # 169 | Philadelphia | Pennsylvania | 24948 | 7123.8 | Philadelphia | 115644.0 | 39.97611 | -75.12472 |

In [12]:

df.columns

Out[12]:

Index(['Zip Code', 'Price', 'Beds', 'Baths', 'Living Space', 'Address', 'City',
'State', 'Zip Code Population', 'Zip Code Density', 'County',
'Median Household Income', 'Latitude', 'Longitude'],
dtype='object')

In [13]:

df.drop(columns={'Address', 'City',
'State', 'County', 'Zip Code'},inplace=True)

In [14]:

df

Out[14]:

| | Price | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|-----------|------|-------|--------------|---------------------|------------------|-------------------------|----------|------------|
| 0 | 3999000.0 | 2 | 3 | 1967 | 29563 | 20967.9 | 370046.0 | 40.72001 | -74.00472 |
| 1 | 3999000.0 | 2 | 3 | 1967 | 29563 | 20967.9 | 370046.0 | 40.72001 | -74.00472 |
| 2 | 1650000.0 | 1 | 1 | 718 | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| 3 | 760000.0 | 3 | 2 | 1538 | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| 4 | 1100000.0 | 1 | 1 | 600 | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 39976 | 2495000.0 | 4 | 4 | 3380 | 22890 | 2086.8 | 205611.0 | 47.65139 | -122.40223 |
| 39977 | 2295000.0 | 4 | 4 | 2878 | 22890 | 2086.8 | 205611.0 | 47.65139 | -122.40223 |
| 39978 | 950000.0 | 3 | 2 | 1380 | 22890 | 2086.8 | 205611.0 | 47.65139 | -122.40223 |
| 39979 | 425000.0 | 2 | 1 | 856 | 22890 | 2086.8 | 205611.0 | 47.65139 | -122.40223 |
| 39980 | 1150000.0 | 3 | 3 | 2840 | 22890 | 2086.8 | 205611.0 | 47.65139 | -122.40223 |

39981 rows × 9 columns

In [15]:

df.isnull().mean()*100

Out[15]:

Price
Beds
Baths
Living Space
Zip Code Population
Zip Code Density
Median Household Income
Latitude
Longitude
dtype: float64

0.000000
0.000000
0.000000
0.000000
0.000000
0.000000
0.005002
0.000000
0.000000
dtype: float64

In [16]:

df.dropna(inplace=True)

```
In [17]: df.head(4)
```

Out[17]:

| | Price | Beds | Baths | Living Space | Zip Code | Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|---|-----------|------|-------|--------------|----------|------------|------------------|-------------------------|----------|-----------|
| 0 | 3999000.0 | 2 | 3 | 1967 | | 29563 | 20967.9 | 370046.0 | 40.72001 | -74.00472 |
| 1 | 3999000.0 | 2 | 3 | 1967 | | 29563 | 20967.9 | 370046.0 | 40.72001 | -74.00472 |
| 2 | 1650000.0 | 1 | 1 | 718 | | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| 3 | 760000.0 | 3 | 2 | 1538 | | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |

```
In [18]: df.duplicated().sum()
```

Out[18]: 1447

```
In [19]: df.drop_duplicates(inplace=True)
```

```
In [20]: df.shape
```

Out[20]: (38532, 9)

```
In [21]: df.head(5)
```

Out[21]:

| | Price | Beds | Baths | Living Space | Zip Code | Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|---|-----------|------|-------|--------------|----------|------------|------------------|-------------------------|----------|-----------|
| 0 | 3999000.0 | 2 | 3 | 1967 | | 29563 | 20967.9 | 370046.0 | 40.72001 | -74.00472 |
| 2 | 1650000.0 | 1 | 1 | 718 | | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| 3 | 760000.0 | 3 | 2 | 1538 | | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| 4 | 1100000.0 | 1 | 1 | 600 | | 29815 | 23740.9 | 249880.0 | 40.73407 | -74.00601 |
| 5 | 764900.0 | 1 | 1 | 643 | | 15514 | 20107.7 | 188289.0 | 40.75235 | -73.97260 |

```
In [22]: df['Baths'].unique()
```

Out[22]: array([3, 1, 2, 10, 4, 5, 9, 6, 11, 8, 7, 12, 24, 21, 20, 14, 15, 28, 66, 18, 16, 17, 36, 19, 37, 42, 46, 56], dtype=int64)

SCALING

```
In [23]: from sklearn.preprocessing import StandardScaler
```

```

In [24]: #Price
Price_scaler=StandardScaler()
df['Price']=Price_scaler.fit_transform(ns.array(df['Price']).reshape(len(df['Price']),1))

#Beds
Beds_scaler=StandardScaler()
df['Beds']=Beds_scaler.fit_transform(ns.array(df['Beds']).reshape(len(df['Beds']),1))

#Baths
Baths_scaler=StandardScaler()
df['Baths']=Baths_scaler.fit_transform(ns.array(df['Baths']).reshape(len(df['Baths']),1))

#Living Space
Living_Space_scaler=StandardScaler()
df['Living Space']=Living_Space_scaler.fit_transform(ns.array(df['Living Space']).reshape(len(df['Living Space']),1))

#Zip Code Population
Zip_Code_Population_scaler=StandardScaler()
df['Zip Code Population']=Zip_Code_Population_scaler.fit_transform(ns.array(df['Zip Code Population']).reshape(len(df['Zip Code P

#Zip Code Density
Zip_Code_Density_scaler=StandardScaler()
df['Zip Code Density']=Zip_Code_Density_scaler.fit_transform(ns.array(df['Zip Code Density']).reshape(len(df['Zip Code Density'])

#Median Household Income
Median_Household_Income_scaler=StandardScaler()
df['Median Household Income']=Median_Household_Income_scaler.fit_transform(ns.array(df['Median Household Income']).reshape(len(df

#Latitude
Latitude_scaler=StandardScaler()
df['Latitude']=Latitude_scaler.fit_transform(ns.array(df['Latitude']).reshape(len(df['Latitude']),1))

#Longitude
Longitude_scaler=StandardScaler()
df['Longitude']=Longitude_scaler.fit_transform(ns.array(df['Longitude']).reshape(len(df['Longitude']),1))

```

```

In [25]: df.head(5)

```

```

Out[25]:

```

| | Price | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|---|----------|-----------|-----------|--------------|---------------------|------------------|-------------------------|----------|-----------|
| 0 | 3.520793 | -0.890824 | 0.403009 | 0.052402 | -0.434679 | 6.338620 | 5.464918 | 0.955387 | 1.596649 |
| 2 | 1.069626 | -1.652084 | -1.099531 | -0.974282 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 3 | 0.140917 | -0.129564 | -0.348261 | -0.300238 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 4 | 0.495705 | -1.652084 | -1.099531 | -1.071279 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 5 | 0.146030 | -1.652084 | -1.099531 | -1.035933 | -1.192575 | 6.045292 | 1.630546 | 0.962657 | 1.598782 |

```

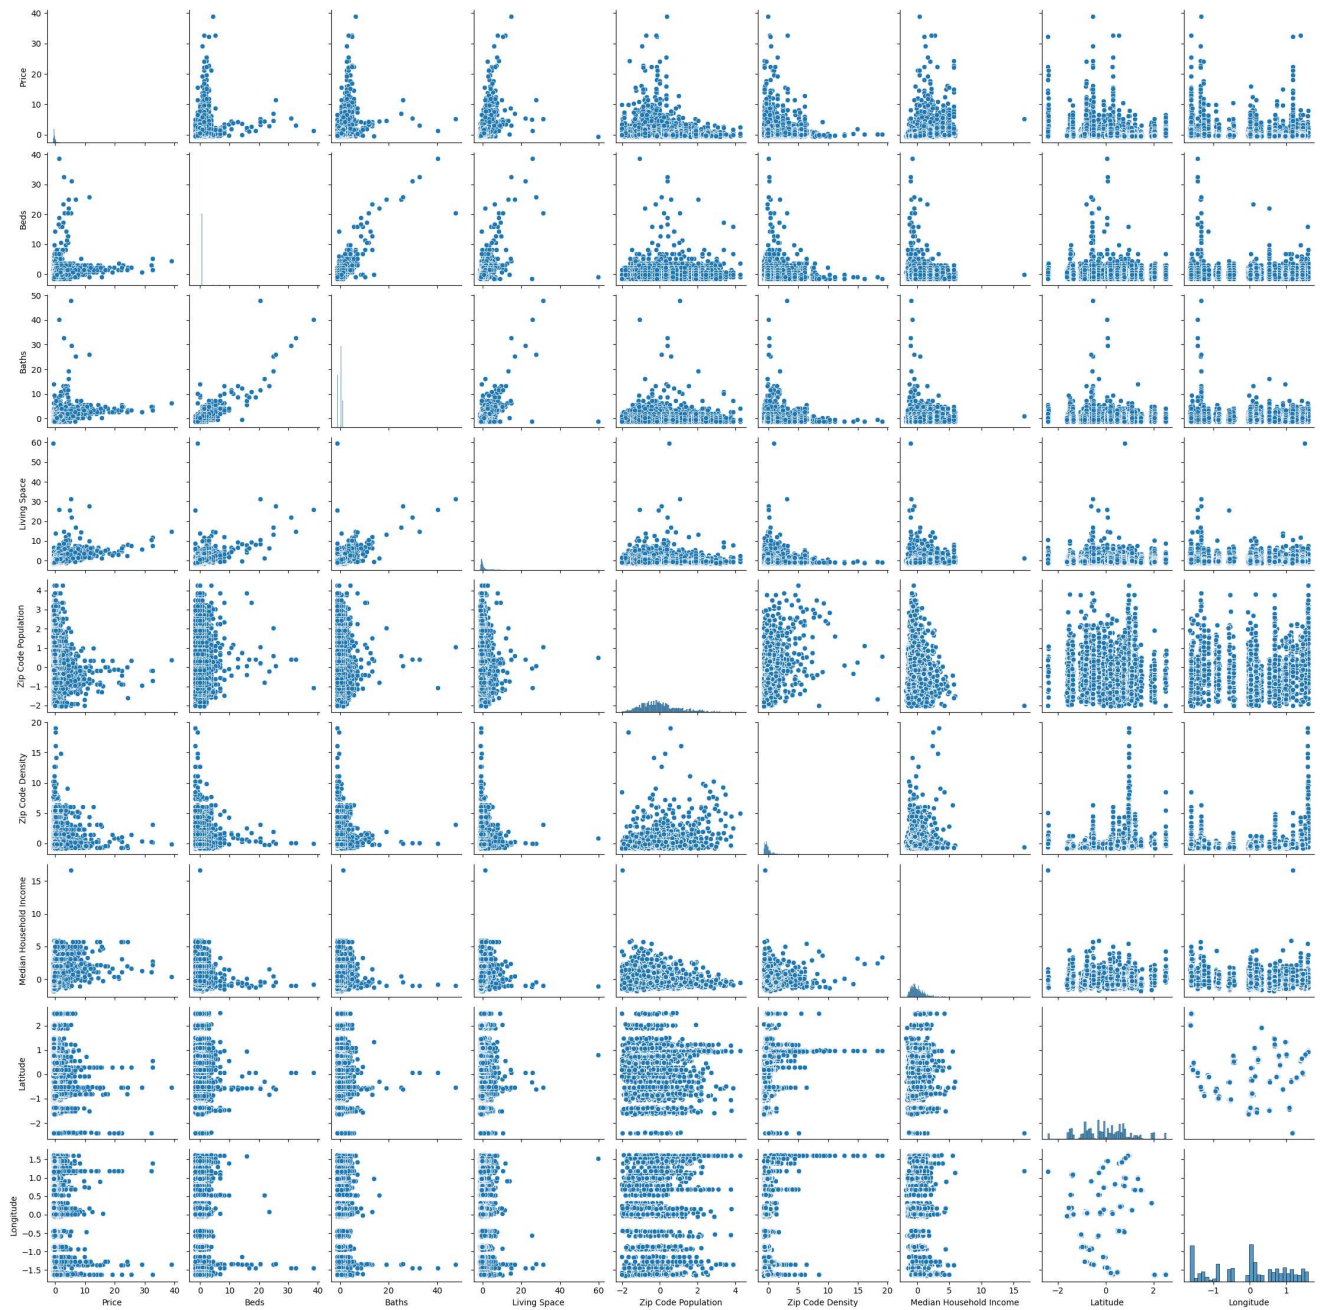
In [26]: import warnings
warnings.filterwarnings('ignore')

```

DATA VISUALIZATION

```
In [27]: sn.pairplot(df)
```

```
Out[27]: <seaborn.axisgrid.PairGrid at 0x264b4d4b110>
```



```
In [28]: sn.set(style='whitegrid')

fig, axs = pl.subplots(3, 2, figsize=(15, 10))

# Price distribution
sn.histplot(df['Price'], bins=30, ax=axs[0, 0], kde=True)
axs[0, 0].set_title('Price Distribution')

# Beds distribution
sn.countplot(x='Beds', data=df, ax=axs[0, 1])
axs[0, 1].set_title('Beds Distribution')

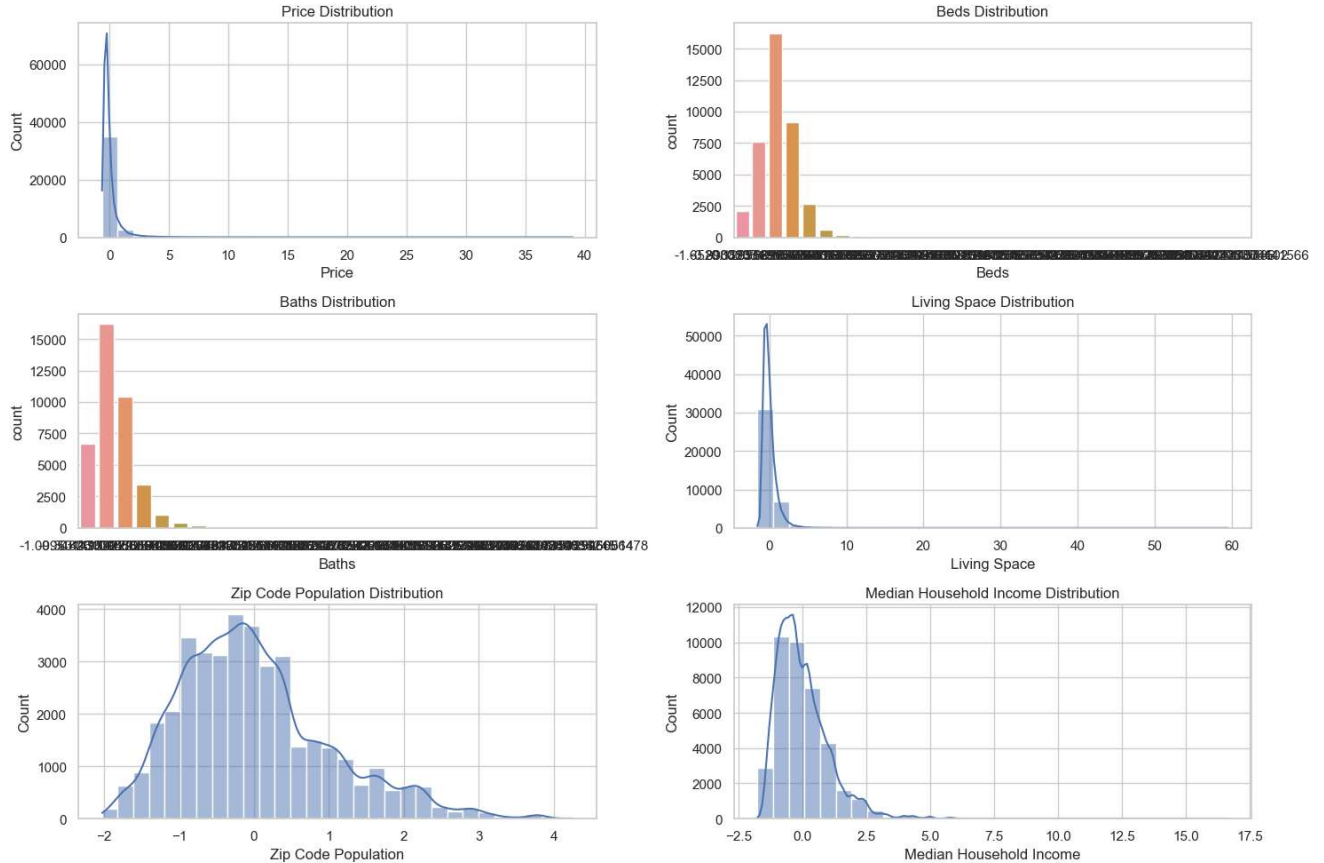
# Baths distribution
sn.countplot(x='Baths', data=df, ax=axs[1, 0])
axs[1, 0].set_title('Baths Distribution')

# Living Space distribution
sn.histplot(df['Living Space'], bins=30, ax=axs[1, 1], kde=True)
axs[1, 1].set_title('Living Space Distribution')

# Zip Code Population distribution
sn.histplot(df['Zip Code Population'], bins=30, ax=axs[2, 0], kde=True)
axs[2, 0].set_title('Zip Code Population Distribution')

# Median Household Income distribution
sn.histplot(df['Median Household Income'], bins=30, ax=axs[2, 1], kde=True)
axs[2, 1].set_title('Median Household Income Distribution')


pl.tight_layout()
pl.show()
```



```
In [29]: plt.figure(figsize=(12, 6))
sn.scatterplot(x='Beds', y='Price', data=df, color='blue', alpha=0.6)
plt.title('Relationship between Beds and Price')
plt.xlabel('Beds')
plt.ylabel('Price')
plt.show()
```



```
In [30]: plt.figure(figsize=(12, 6))
sn.scatterplot(x='Baths', y='Price', data=df, color='green', alpha=0.6)
plt.title('Relationship between Baths and Price')
plt.xlabel('Baths')
plt.ylabel('Price')
plt.show()
```




```
In [31]: pl.figure(figsize=(12, 6))
sn.scatterplot(x='Living Space', y='Price', data=df, color='orange', alpha=0.6)
pl.title('Relationship between Living Space and Price')
pl.xlabel('Living Space (sqft)')
pl.ylabel('Price')
pl.show()
```



```
In [32]: df
```


Out[32]:

| | Price | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|-----------|-----------|-----------|--------------|---------------------|------------------|-------------------------|----------|-----------|
| 0 | 3.520793 | -0.890824 | 0.403009 | 0.052402 | -0.434679 | 6.338620 | 5.464918 | 0.955387 | 1.596649 |
| 2 | 1.069626 | -1.652084 | -1.099531 | -0.974282 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 3 | 0.140917 | -0.129564 | -0.348261 | -0.300238 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 4 | 0.495705 | -1.652084 | -1.099531 | -1.071279 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 5 | 0.146030 | -1.652084 | -1.099531 | -1.035933 | -1.192575 | 6.045292 | 1.630546 | 0.962657 | 1.598782 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 39976 | 1.951378 | 0.631697 | 1.154280 | 1.213896 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39977 | 1.742680 | 0.631697 | 1.154280 | 0.801249 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39978 | 0.339181 | -0.129564 | -0.348261 | -0.430115 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39979 | -0.208653 | -0.890824 | -1.099531 | -0.860846 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39980 | 0.547880 | -0.129564 | 0.403009 | 0.770013 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |

38532 rows × 9 columns


```
In [33]: Median Household Income, have a high correlation with price
'Beds')[['Baths', 'Price']].mean().sort_values('Price', ascending=False).reset_index().head(10), x='Baths', y='Price', hue='Beds')
```

```
Out[33]: <seaborn.axisgrid.FacetGrid at 0x264ccb2d210>
```



```
In [34]: sn.scatterplot(x='Baths',y='Price',data=df)
```

```
Out[34]: <Axes: xlabel='Baths', ylabel='Price'>
```



LINEAR REGRESSION MODEL

```
In [35]: X = df.drop('Price',axis=1)
Y= df['Price']
```

```
In [36]: X
```

Out[36]:

| | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|-----------|-----------|--------------|---------------------|------------------|-------------------------|----------|-----------|
| 0 | -0.890824 | 0.403009 | 0.052402 | -0.434679 | 6.338620 | 5.464918 | 0.955387 | 1.596649 |
| 2 | -1.652084 | -1.099531 | -0.974282 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 3 | -0.129564 | -0.348261 | -0.300238 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 4 | -1.652084 | -1.099531 | -1.071279 | -0.421084 | 7.284213 | 2.929879 | 0.958548 | 1.596564 |
| 5 | -1.652084 | -1.099531 | -1.035933 | -1.192575 | 6.045292 | 1.630546 | 0.962657 | 1.598782 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 39976 | 0.631697 | 1.154280 | 1.213896 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39977 | 0.631697 | 1.154280 | 0.801249 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39978 | -0.129564 | -0.348261 | -0.430115 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39979 | -0.890824 | -1.099531 | -0.860846 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |
| 39980 | -0.129564 | 0.403009 | 0.770013 | -0.794665 | -0.099833 | 1.995973 | 2.513491 | -1.617405 |

38532 rows × 8 columns

```
In [37]: Y
```

Out[37]:

| | |
|-------|-----------|
| 0 | 3.520793 |
| 2 | 1.069626 |
| 3 | 0.140917 |
| 4 | 0.495705 |
| 5 | 0.146030 |
| ... | ... |
| 39976 | 1.951378 |
| 39977 | 1.742680 |
| 39978 | 0.339181 |
| 39979 | -0.208653 |
| 39980 | 0.547880 |

Name: Price, Length: 38532, dtype: float64

```
In [38]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [39]: X_train, X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=42)
```

```
In [40]: model = LinearRegression()
```

```
In [41]: model.fit(X_train,Y_train)
```

Out[41]:

| |
|--------------------|
| LinearRegression |
| LinearRegression() |

```
In [42]: y_pred = model.predict(X_test)
y_pred
```

Out[42]:

| |
|---|
| array([-0.60675959, 1.76218257, 0.3105298 , ..., -0.34763035, 0.1734936 , -0.42821048]) |
|---|

```
In [43]: X_test
```

Out[43]:

| | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|-----------|-----------|--------------|---------------------|------------------|-------------------------|-----------|-----------|
| 13412 | -0.129564 | -0.348261 | -0.286264 | -0.335687 | 0.085228 | -1.193953 | 1.487314 | 0.668800 |
| 36437 | 0.631697 | 1.905550 | 2.290722 | -0.295443 | 1.212198 | 1.327458 | 0.301761 | -1.606663 |
| 10219 | -0.129564 | 1.154280 | 0.590816 | -0.748864 | -0.552169 | 0.407560 | 0.402746 | 0.830379 |
| 37059 | -0.129564 | 1.154280 | 0.032674 | 0.855454 | 0.155098 | 2.400007 | 0.185601 | -1.583659 |
| 527 | -0.890824 | -0.348261 | -0.167073 | 0.507013 | 5.745996 | -0.063854 | 0.956839 | 1.606164 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 27479 | -0.129564 | 0.403009 | 0.015412 | 0.073768 | -0.398003 | 0.377709 | 0.558808 | -0.446584 |
| 36363 | 0.631697 | 0.403009 | 0.265301 | -0.208535 | 0.803646 | 0.764233 | 0.307116 | -1.609700 |
| 25241 | -0.129564 | 0.403009 | -0.335584 | -0.845212 | -0.371303 | -1.582460 | -1.057238 | -0.556380 |
| 29606 | -0.129564 | -0.348261 | 0.478201 | -0.163112 | -0.561206 | -0.115772 | -0.928613 | -0.859181 |
| 31296 | -0.890824 | -0.348261 | -0.708775 | 0.352887 | -0.609935 | -0.659103 | -0.316463 | -0.559348 |

7707 rows × 8 columns

In [44]:

Y_test

Out[44]:

13412 -0.422569
36437 0.901624
10219 0.015698
37059 0.495588
527 0.326659
...
27479 -0.134878
36363 0.390312
25241 -0.479961
29606 0.014654
31296 -0.433004
Name: Price, Length: 7707, dtype: float64

In [45]:

X_train

Out[45]:

| | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|-----------|-----------|--------------|---------------------|------------------|-------------------------|-----------|-----------|
| 37751 | 0.631697 | -0.348261 | -0.394769 | -0.368864 | 0.204884 | -0.895359 | 0.460276 | -1.553601 |
| 29600 | -0.129564 | -0.348261 | -0.241876 | -0.163112 | -0.561206 | -0.115772 | -0.928613 | -0.859181 |
| 25684 | -0.129564 | 0.403009 | 0.613010 | -0.486953 | -0.537063 | -0.116130 | -1.032313 | -0.568438 |
| 26492 | 0.631697 | -0.348261 | -0.249274 | 1.606121 | 0.390627 | -0.683828 | 0.724966 | -0.464011 |
| 2632 | 1.392957 | -0.348261 | 0.285030 | -0.921925 | 0.218013 | -1.242770 | 0.632634 | 1.420763 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 6634 | -0.129564 | -0.348261 | -0.176115 | -0.105335 | -0.238279 | -0.906624 | -1.379561 | 1.093294 |
| 11769 | -0.129564 | -1.099531 | -0.196665 | 0.023328 | -0.357629 | -1.208383 | 0.746528 | 0.785786 |
| 39575 | -0.890824 | -0.348261 | -0.599448 | 0.610482 | 0.220980 | 0.716091 | 2.488980 | -1.608539 |
| 932 | -0.890824 | -1.099531 | -0.750697 | 1.185175 | 1.465288 | -1.213826 | 0.791183 | 1.523356 |
| 16356 | -0.129564 | -1.099531 | -0.203242 | 0.351053 | -0.347467 | -0.692414 | 0.277075 | 0.040551 |

30825 rows × 8 columns

In [46]:

X_test

Out[46]:

| | Beds | Baths | Living Space | Zip Code Population | Zip Code Density | Median Household Income | Latitude | Longitude |
|-------|-----------|-----------|--------------|---------------------|------------------|-------------------------|-----------|-----------|
| 13412 | -0.129564 | -0.348261 | -0.286264 | -0.335687 | 0.085228 | -1.193953 | 1.487314 | 0.668800 |
| 36437 | 0.631697 | 1.905550 | 2.290722 | -0.295443 | 1.212198 | 1.327458 | 0.301761 | -1.606663 |
| 10219 | -0.129564 | 1.154280 | 0.590816 | -0.748864 | -0.552169 | 0.407560 | 0.402746 | 0.830379 |
| 37059 | -0.129564 | 1.154280 | 0.032674 | 0.855454 | 0.155098 | 2.400007 | 0.185601 | -1.583659 |
| 527 | -0.890824 | -0.348261 | -0.167073 | 0.507013 | 5.745996 | -0.063854 | 0.956839 | 1.606164 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 27479 | -0.129564 | 0.403009 | 0.015412 | 0.073768 | -0.398003 | 0.377709 | 0.558808 | -0.446584 |
| 36363 | 0.631697 | 0.403009 | 0.265301 | -0.208535 | 0.803646 | 0.764233 | 0.307116 | -1.609700 |
| 25241 | -0.129564 | 0.403009 | -0.335584 | -0.845212 | -0.371303 | -1.582460 | -1.057238 | -0.556380 |
| 29606 | -0.129564 | -0.348261 | 0.478201 | -0.163112 | -0.561206 | -0.115772 | -0.928613 | -0.859181 |
| 31296 | -0.890824 | -0.348261 | -0.708775 | 0.352887 | -0.609935 | -0.659103 | -0.316463 | -0.559348 |

7707 rows × 8 columns

In [47]:

mae=mean_absolute_error(Y_test,y_pred)
mse = mean_squared_error(Y_test, y_pred)
rmse= ns.sqrt(mse)
r2 =r2_score(Y_test,y_pred)

In [48]:

print('Mean Absolute Error',mae)
print('Mean Squared error',mse)
print('root Mean Squared error',rmse)
print('R2 score',r2)

Mean Absolute Error 0.32099721054789543
Mean Squared error 0.7689516707371317
root Mean Squared error 0.8768988942501477
R2 score 0.4134816287003974

In [49]:

model.coef_

Out[49]:

array([-0.07371379, 0.15864634, 0.3635581 , -0.01728186, 0.15404179,
0.24577778, -0.06971322, -0.11222055])

```
In [50]: model.intercept_
```

```
Out[50]: -0.0037304625161884517
```